

Item-Based Collaborative Filtering for Music Recommendation

A Co-occurrence Matrix Approach

Your Name

June 29, 2025

Overview

- 1 Problem Statement & Dataset
- 2 Methodology
- 3 Technical Implementation
- 4 Evaluation & Results
- 5 Advantages & Limitations
- 6 Conclusion

Music Recommendation Challenge

Objective

Build a personalized music recommendation system using item-based collaborative filtering

Dataset Structure:

- **user_id**: Unique identifier for each user
- **song_id**: Unique identifier for each track
- **listen_count**: Number of times user played the song

Data Characteristics:

- No ratings, only interaction counts
- Real-world music streaming data

Key Statistics:

- Unique users: Large scale
- Unique songs: Thousands
- Train/Test split: 80/20
- Data format: User-Item interaction matrix

Data Quality:

- No missing values
- No duplicate entries

Dataset Details & Preprocessing

Core Data Fields

user_id User identifier - represents individual music listeners

song_id Track identifier - represents individual songs in the catalog

listen_count Frequency of interaction - how many times user played the song

Why These Fields Matter:

- **user_id**: Enables personalization
- **song_id**: Items to recommend
- **listen_count**: Measures preference strength

Data Insights:

- Co-occurrence patterns reveal similarity

Preprocessing Steps:


- 1 Load user-track interaction data
- 2 Filter users with 30 unique songs
- 3 Split into train/test (80/20)
- 4 Create user-item interaction matrix

Result: Clean, dense dataset focusing on engaged users with sufficient listening history

Item Similarity Principle

Core Assumption

"Songs listened to by similar users are likely to be similar to each other"



Algorithm 1 Item Similarity Recommender

- 1: **Input:** User ID, Training Data
 - 2: Get user's listening history: $S_u = \{s_1, s_2, \dots, s_k\}$
 - 3: Get all songs in system: $S_{all} = \{s_1, s_2, \dots, s_n\}$
 - 4: **For each** song pair (s_i, s_j) :
 - 5: Compute Jaccard similarity using co-occurrence
 - 6: Build co-occurrence matrix $M_{k \times n}$
 - 7: Compute weighted similarity scores
 - 8: Rank songs by similarity scores
 - 9: Filter out already-listened songs
 - 10: **Return:** Top-10 recommendations
-

Co-occurrence Matrix Construction

Jaccard Similarity Formula

$$J(s_i, s_j) = \frac{|U_{s_i} \cap U_{s_j}|}{|U_{s_i} \cup U_{s_j}|} \quad (1)$$

where U_{s_i} = set of users who listened to song s_i

Matrix Dimensions: $k \times n$

- k = user's songs, n = total songs
- Each cell = Jaccard similarity

Quick Example: Song A: {User1, User2, User3}, Song B: {User2, User3, User4}

Jaccard = $\frac{2}{4} = 0.5$ (50% similar)

Matrix: $k \times n$ (k = user's songs, n = total)

Weighted Score Calculation

$$\text{Score}(s_j) = \frac{1}{k} \sum_{i=1}^k J(s_i, s_j) \quad (2)$$

Example: User {Rock, Pop, Jazz},
Blues similarities {0.6, 0.2, 0.8}

$$\text{Score} = \frac{0.6+0.2+0.8}{3} = 0.53$$

Key Implementation Details

Data Preprocessing

- Filter users with 30 unique songs (active users)
- 80/20 train-test split

Similarity Computation

```
# Jaccard Index calculation
users_intersection = users_i.intersection(users_j)
users_union = users_i.union(users_j)
similarity = len(users_intersection) / len(users_union)
```

Recommendation Generation

- Compute average similarity across all user's songs
- Sort by similarity score (descending)
- Filter out previously listened songs
- Return top-10 recommendations

Evaluation Methodology

Metrics Used

- **Precision@K:** $\frac{\text{Relevant items in top-K}}{K}$
- **Recall@K:** $\frac{\text{Relevant items in top-K}}{\text{Total relevant items}}$
- Evaluated for $K = 1, 2, \dots, 10$

Experimental Setup

- User sample: 5% of test users
- Test data: 1000 records subset
- Recommendations: Top-10 per user
- Ground truth: User's actual listening in test set

Precision-Recall Curve: Shows trade-off between precision and recall across different K values

Results: Impact of User Filtering

Before Filtering (All Users):

- **Precision@10:** Up to 20%
- **Recall@10:** Up to 12.5%
- Curve position: Center to upper-left
- Includes casual users (1-5 songs)

After Filtering (30 songs):

- **Precision@10:** 7.5-10%
- **Recall@10:** 2.5%
- Curve position: Lower-left corner
- Only active, engaged users

Performance Comparison:


Metric	Before	After
Max Precision	20%	10%
Max Recall	12.5%	2.5%
User Quality	Mixed	High
Realism	Low	High

Why the Drop?

- Casual users = easy predictions
- Active users = complex preferences
- Filtered results more realistic


Precision-Recall Curves Comparison

Before Filtering



item_similarity_model

After Filtering (30 songs)



item_similarity_model

Key Insight

Performance drop is expected and desirable - it indicates proper evaluation methodology focusing on meaningful user segments

Why User Filtering Improves Evaluation Quality

The "Easy Prediction" Problem

Casual Users (1-5 songs): Small test sets → Artificially high precision/recall

Mathematical Example:

Casual User:

Test songs = 1

Correct predictions = 1

$$\text{Precision} = \frac{1}{1} = 100\%$$

Active User:

Test songs = 8

Correct predictions = 1

$$\text{Precision} = \frac{1}{8} = 12.5\%$$

Business Relevance:

- **Active users:** Generate revenue, stay engaged
- **Casual users:** May not return, limited value
- **Filtering focus:** Real users
- **Realistic metrics:** True system performance

Industry Standard: Music platforms evaluate on users with sufficient listening history (20-50 songs)

Strengths & Weaknesses

Advantages:

- **Personalized:** Based on actual user behavior
- **Explainable:** "Because you liked X, try Y"
- **No content needed:** Only interaction data
- **Domain agnostic:** Works across different domains
- **Quality recommendations:** Similar items tend to be relevant

Limitations:

- **Cold start:** Can't recommend to new users
- **Sparsity:** Matrix mostly zeros
- **Scalability:** Expensive computation for large datasets
- **Echo chamber:** May lack diversity
- **Popular bias:** Tends to recommend popular items

Best Use Cases

Music streaming, e-commerce, movie recommendations where user-item interactions are abundant

Summary & Future Work

Key Achievements

- Successfully implemented item-based collaborative filtering
- **Demonstrated importance of user filtering** for realistic evaluation
- Achieved meaningful precision-recall performance (7-10% precision for active users)
- **Revealed evaluation bias** in unfiltered data (artificially high metrics)
- Created scalable, interpretable recommendation system

Critical Insights

- **User filtering is essential:** Prevents inflated metrics from casual users
- **Performance drop is good:** Indicates realistic, challenging evaluation
- **Focus on engaged users:** More business-relevant metrics

Thank you for your attention!